Chest X-Ray Pneumonia Detection Using Deep Learning

School of Computer Engineering

(Manipal Institute of Technology, Manipal, India)

Abstract: Medical imaging plays a pivotal role in the early and accurate diagnosis of critical diseases such as pneumonia and brain tumors. Pneumonia is a serious respiratory disease that affects millions of people worldwide and requires timely diagnosis for effective treatment. Chest X-ray imaging is widely used for detecting pneumonia; however, manual interpretation can be time-consuming and prone to errors. This project proposes an automated system using deep learning techniques to classify chest X-ray images into Pneumonia and Normal categories. A Convolutional Neural Network (CNN) with transfer learning is utilized to extract meaningful features and improve classification accuracy. The system incorporates preprocessing techniques such as normalization, noise reduction, and data augmentation to enhance model performance. Preliminary results indicate promising accuracy, demonstrating the potential of AI-assisted diagnosis in improving healthcare efficiency and reducing workload on radiologists. This work provides actionable insights for deploying AI-driven diagnostic systems in real-world healthcare settings.

I. Introduction

Pneumonia is a serious respiratory infection that affects the lungs and remains one of the leading causes of mortality worldwide, particularly among children and elderly individuals. It is characterized by inflammation of the alveoli, where fluid accumulation can lead to symptoms such as cough, breathing difficulty, and reduced oxygen intake. In severe cases, pneumonia can result in respiratory failure and death. Due to its widespread impact and high mortality rate, early and accurate diagnosis is essential for effective treatment and improved patient outcomes.

Chest X-ray (CXR) imaging is the most commonly used diagnostic method for detecting pneumonia because of its low cost, accessibility, and efficiency in identifying lung abnormalities. However, accurate interpretation of chest X-rays requires significant expertise and experience. In real-world clinical settings, radiologists often face challenges such as high patient volumes, fatigue, and variability in judgment, which can lead to delays or errors in diagnosis. This problem is further intensified in regions with limited access to skilled healthcare professionals, creating a strong need for reliable automated diagnostic systems .

Computer-Aided Detection (CAD) systems have emerged as a promising solution to address these challenges by enhancing diagnostic accuracy, reducing inter-observer variability, and enabling faster analysis of medical images. With advancements in artificial intelligence, particularly in machine learning (ML) and deep learning (DL), it has become possible to automatically analyze complex medical imaging data with high precision. Among these techniques, Convolutional Neural Networks (CNNs) have demonstrated exceptional performance in image classification tasks due to their ability to learn hierarchical features directly from raw images.

In this project, a deep learning-based approach is proposed to automatically detect pneumonia from chest X-ray images. By leveraging CNNs along with transfer learning and the effective preprocessing techniques, the system aims to accurately classify images into Pneumonia and Normal categories(binary). The objective is to support the healthcare professionals in early diagnosis, improve screening efficiency, and reduce dependence on manual interpretation, particularly in the resource-constrained environments.

II. Methodology

The proposed system follows a structured approach for detecting pneumonia from chest X-ray images using both machine learning and deep learning techniques.

Initially, multiple approaches were explored, including traditional machine learning models such as Support Vector Machine (SVM), standalone deep learning models using Convolutional Neural Networks (CNNs), and hybrid models combining CNN with XGBoost. This comparative analysis was performed to identify the most effective model for the task.
The dataset used consists of labeled chest X-ray images categorized into Normal and Pneumonia classes. Preprocessing steps such as image resizing, normalization, grayscale conversion, and data augmentation were applied to improve model performance and reduce overfitting.

For model development, CNN-based architectures were used for automatic feature extraction. Transfer learning techniques using pre-trained models such as VGG16 were also implemented to enhance performance on limited medical data. In addition, traditional machine learning approaches were tested using extracted

features.

Based on preliminary experimental results, the hybrid CNN + XGBoost model demonstrated superior performance in terms of accuracy, stability, and generalization compared to standalone models. Therefore, further development and optimization in this project focus on refining the CNN + XGBoost approach.

III. Dataset

A. Chest X-Ray Dataset:

The dataset used in this project is the Chest X-Ray Images (Pneumonia) dataset obtained from Kaggle. It consists of labeled chest X-ray images categorized into two classes: Pneumonia and Normal. This dataset is widely used for medical image classification tasks and provides a suitable benchmark for evaluating deep learning models.

The dataset contains total of 5856 images, distributed as:

- Pneumonia: 3875 images
- Normal: 1341 images

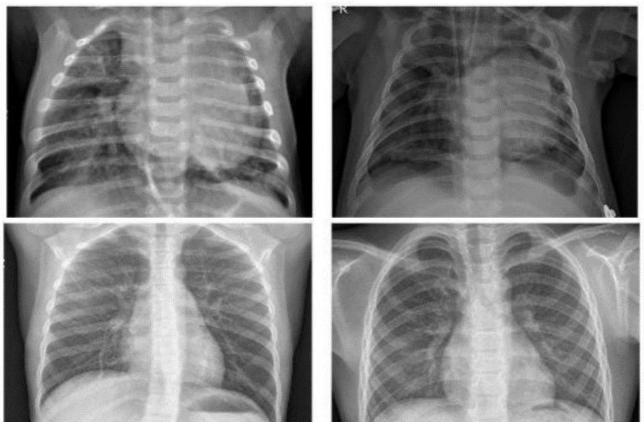

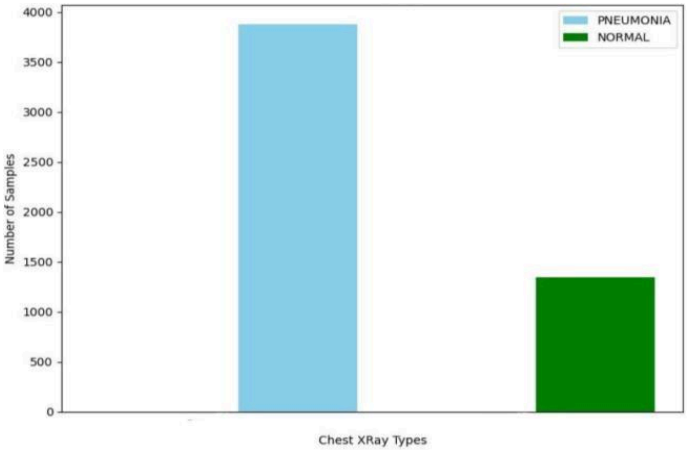

This class imbalance reflects real-world scenarios where disease cases may be more prevalent in clinical datasets. A histogram representation of the dataset distribution is shown in Figure 2, highlighting the difference between the two classes.

All images are stored in grayscale format, as chest X-rays inherently capture variations in radiation absorption. The lungs appear darker, while bones and dense structures appear brighter. The original images are high resolution (approximately 2000 × 2000 pixels) and are resized during preprocessing for computational efficiency.

B. Dataset Characteristics:

- Image Type: Grayscale chest X-ray images
- Format: JPEG
- Resolution: High-resolution images (resized during preprocessing)
- Classes: Binary classification (Pneumonia/Normal)

The dataset also includes visually distinct patterns:

- Pneumonia images show opacities or white patches in lung regions
- Normal images show clear lung fields with minimal obstruction

Sample images for both classes are illustrated in Figure 1, demonstrating the visual differences learned by the model.

C. Ethical Considerations:

Although this project uses a publicly available dataset, ethical guidelines for handling medical data are important:

- Informed Consent: Data should be collected with proper patient consent.
- Anonymization: All personal identifiers must be removed to ensure privacy.
- Regulatory Compliance: Standards such as HIPAA and GDPR should be followed.
- Data Security: Proper measures should be taken to prevent unauthorized access.
- Transparency: Data usage should be clearly defined and limited to research purposes.
- Accountability: Ethical oversight should be maintained throughout the project.

These principles ensure responsible use of medical data and are essential for real-world healthcare applications.

IV. System Design

The proposed system follows a hybrid architecture combining deep learning and machine learning techniques for efficient pneumonia detection.

The workflow of the system is illustrated in the hybrid model diagram (CNN + XGBoost pipeline), where the process is divided into multiple stages:

- ***Data Loading and Preprocessing:***
  Input chest X-ray images are resized, normalized, and enhanced. Data augmentation techniques are applied to improve generalization.
- ***Dataset Splitting:***
  The dataset is divided into training, validation, and testing sets. Different split ratios were experimented with to identify optimal performance.
- ***Feature Extraction using CNN:***
  A pre-trained CNN model (such as VGG16) is used to extract high-level feature representations from images. Transfer learning allows the model to utilize previously learned features, reducing the need for large datasets.
- ***Feature Vector Conversion:***
  The extracted feature maps are flattened into one-dimensional vectors suitable for machine learning models.
- ***Classification using XGBoost:***
  The feature vectors are passed into an XGBoost classifier, which performs final classification into Pneumonia or Normal categories.
- ***Prediction and Evaluation:***
  The system outputs predictions and evaluates performance using metrics such as accuracy, F1-score, and log loss.

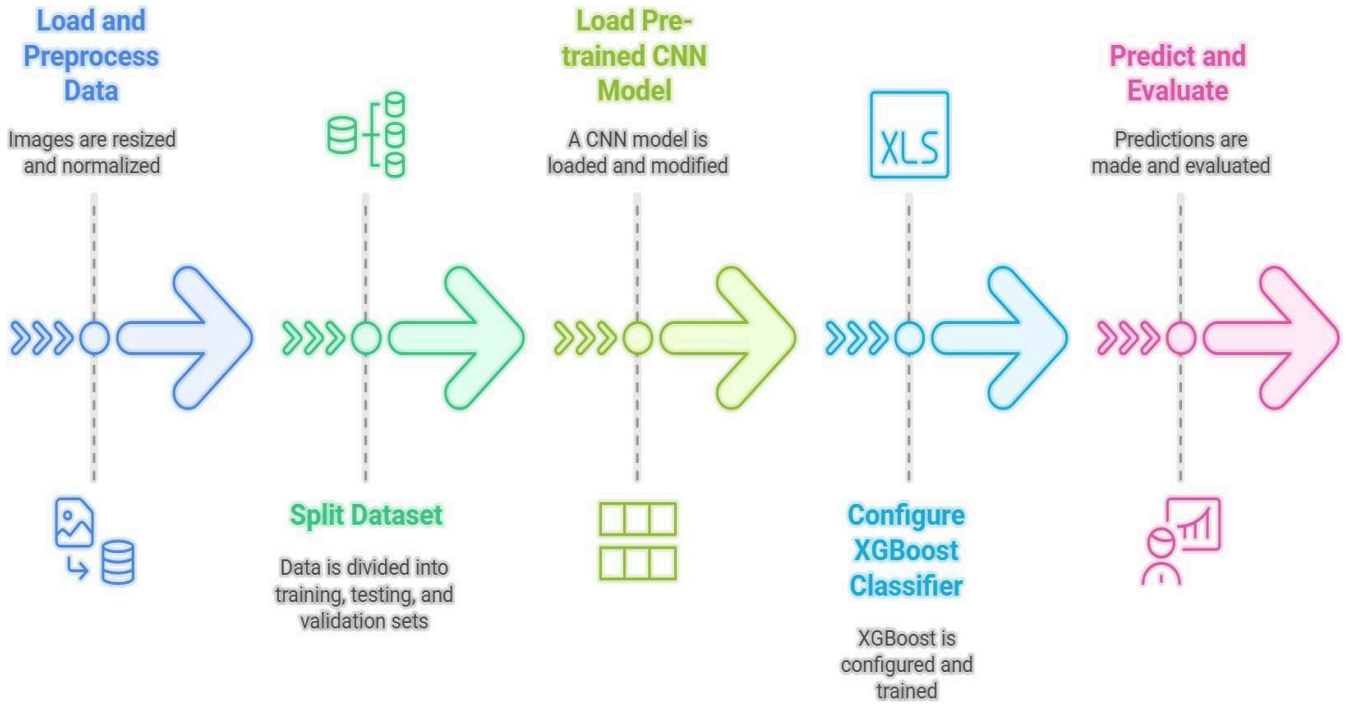

V. Results

The performance of different models was evaluated using multiple training splits and evaluation metrics.

A. SVM Performance

SVM models achieved good accuracy at higher training splits (80/20), reaching approximately 97% accuracy, but their performance declined as the training data decreased. This indicates limited generalization capability for complex image data.

B. CNN Performance

CNN models using transfer learning showed improved performance compared to traditional ML approaches.

Best performance achieved at 60/5/35 split

- Validation Accuracy: 92.01%
- Validation Loss: 0.2316
- ROC-AUC: ~0.90

The results indicate that CNNs are effective in capturing important image features for pneumonia detection.

However, training instability was observed:

- Smooth decrease in training loss
- Erratic validation loss with spikes, indicating possible overfitting or learning rate issues

C. XGBoost and Hybrid Model Performance

Among all approaches, the hybrid CNN + XGBoost model achieved the best results:

- Lowest validation log loss: 0.106
- Faster convergence (fewer boosting rounds ~72)
- More stable training behavior

Compared to:

- Raw XGBoost (higher log loss ~0.146)
- XGBoost + PCA (log loss ~0.109)

The hybrid model consistently outperformed all other approaches.

VI. Preliminary / Expected Outcomes

Based on experimental observations the hybrid CNN + XGBoost model is expected to provide:

- Higher accuracy and better generalization
- Less overfitting compared to standalone CNN
- Stable and reliable predictions
- The system can be used as:
- A clinical decision support tool
- A screening system in resource-limited settings

Further improvements are expected through hyperparameter tuning and larger datasets.

VII. Discussion

The experiments show a clear progression in performance from traditional machine learning models to deep learning and finally hybrid approaches. While SVM performs well on simpler patterns, it lacks the ability to capture complex spatial features in medical images. CNNs improve performance significantly by learning hierarchical representations directly from data.

Impact of Transfer Learning

Transfer learning plays a crucial role in improving CNN performance. Since medical datasets are relatively small, using pre-trained models helps extract meaningful features without requiring extensive training. This is particularly effective for binary classification tasks like pneumonia detection, where patterns are more distinct.

Training Stability and Limitations of CNN

Although CNNs achieved high accuracy, training instability was observed. The validation loss showed irregular spikes despite decreasing training loss, indicating potential issues such as:

- Overfitting
- Improper learning rate
- Sensitivity to dataset split

This highlights that high accuracy alone does not guarantee reliable performance in real-world applications.

Advantages of Hybrid CNN + XGBoost Model

The hybrid model addresses many limitations of standalone CNNs:

Improved Stability: XGBoost provides smoother and more consistent learning compared to CNN classifiers.

Better Generalization: CNN extracts strong features, while XGBoost effectively separates classes using these features.

Lower Log Loss: Reduced prediction uncertainty makes the model more reliable for medical diagnosis.

Efficient Training: Fewer boosting rounds indicate faster convergence and better feature quality.

Data Efficiency Observation

An interesting observation is that binary classification tasks like chest X-ray pneumonia detection can achieve high accuracy even with moderate training data. This is because the class boundaries are relatively well-defined compared to multi-class problems.

Clinical Relevance

From a practical perspective, the hybrid model offers: More stable predictions; Reduced variability; Better suitability for deployment in healthcare systems.

Method Type:-	**Example Algorithms**	**Feature Extraction**	**Interpretability**
Isolated ML	SVM, XGBoost	Manual	High
Isolated DL	CNN	Automatic	Low
Hybrid	CNN + SVM/XGBoost	Automatic + ML	Moderate

VIII. Conclusion

In this project, a deep learning-based system was developed to detect pneumonia from chest X-ray images. The approach uses convolutional neural networks along with transfer learning to classify images into normal and pneumonia categories. Image preprocessing techniques such as resizing, normalization, and augmentation were applied to improve the quality of input data and reduce overfitting.

The results obtained show that the model is capable of identifying pneumonia cases with good accuracy and consistency. Compared to traditional manual analysis, the automated system provides faster predictions and can support doctors in making timely decisions. This is especially useful in situations where there is a shortage of experienced radiologists or a high volume of patients.

Despite these advantages, the system has some limitations. The performance of the model depends on the quality and size of the dataset, and it may not perform equally well on completely new or diverse data. Additionally, the current model only performs binary classification and does not provide detailed medical insights.

In future work, the system can be improved by training on larger datasets, extending it to detect multiple lung diseases, and deploying it as a web or mobile application. With further development, such AI-based tools can become an important support system in healthcare for faster and more efficient diagnosis

IX. References

[1] P. Rajpurkar, J. Irvin, K. Zhu, et al., “CheXNet: Radiologist-Level Pneumonia Detection on Chest X-Rays with Deep Learning,” arXiv preprint arXiv:1711.05225, 2017.

[2] X. Wang, Y. Peng, L. Lu, Z. Lu, M. Bagheri, and R. M. Summers, “ChestX-ray8: Hospital-scale chest X-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases,” in Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR), 2017, pp. 2097–2106.

[3] P. Szepesi and L. Szilágyi, “Detection of pneumonia using convolutional neural networks and deep learning,” Biocybernetics and Biomedical Engineering, vol. 42, no. 3, pp. 1011–1022, 2022.

[4] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining, 2016, pp. 785–794.

[5] Y. Hedhoud, T. Mekhaznia, and M. Amroune, “An improvement of the CNN-XGBoost model for pneumonia disease classification,” Polish Journal of Radiology, vol. 88, pp. 483–493, 2023